

Table 6: Results of instruction-tuning the **Gemma-7B** model with the **Orca-math** dataset.

Epochs	Methods	Average (%)	ARC (%) (norm)	HellaSwag (%) (norm)	MMLU (%)	TrustfulQA (%) (mc2)	Winogrande (%)	GSM8K (%) (strict-match)
1	random	66.80	60.84	81.71	61.51	52.21	77.11	67.40
2	random	66.95	61.35	81.89	60.52	53.44	75.22	69.29
	attention	67.54	62.37	81.93	60.66	55.83	76.87	67.55
	loss	67.01	61.43	82.11	61.00	55.23	76.95	65.35
	length	66.87	61.86	81.64	60.10	53.86	76.09	67.70
3	random	67.09	61.35	81.94	60.56	53.09	75.45	70.13
	attention	66.64	62.20	82.03	60.18	54.74	76.87	63.84
	loss	66.96	61.69	82.08	60.63	55.42	77.27	64.67
	length	66.25	62.12	81.66	60.52	54.37	75.93	62.93
Base Model		63.30	61.00	82.54	62.82	45.07	77.74	50.64

Table 7: Results of instruction-tuning the **Gemma-7B** model with the **SlimOrca-Dedup** dataset.

Epochs	Methods	Average (%)	ARC (%) (norm)	HellaSwag (%) (norm)	MMLU (%)	TrustfulQA (%) (mc2)	Winogrande (%)	GSM8K (%) (strict-match)
1	random	66.49	62.97	81.85	61.76	54.40	76.48	61.49
2	random	66.50	64.25	82.07	60.75	53.31	75.85	62.77
	attention	66.73	64.08	81.37	62.11	53.24	74.82	64.75
	loss	64.89	63.40	81.43	61.27	51.90	75.37	55.95
	length	66.79	64.16	81.89	62.26	53.63	76.56	62.24
3	random	66.36	63.14	81.90	61.19	55.88	74.66	61.41
	attention	66.87	63.65	81.59	62.30	54.98	76.56	62.17
	loss	65.48	64.08	81.30	61.87	52.93	75.85	56.86
	length	66.55	62.54	81.46	61.80	55.98	74.98	62.55
Base Model		63.30	61.00	82.54	62.82	45.07	77.74	50.64

based on gemini research and technology. *arXiv preprint arXiv:2403.08295*.

Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. *Advances in neural information processing systems*, 30.

Yizhong Wang, Yeganeh Kordi, Swaroop Mishra, Alisa Liu, Noah A Smith, Daniel Khashabi, and Hananeh Hajishirzi. 2022. Self-instruct: Aligning language models with self-generated instructions. *arXiv preprint arXiv:2212.10560*.

Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. 2019. Hellaswag: Can a machine really finish your sentence? *arXiv preprint arXiv:1905.07830*.